

# SQL: Schema Definition, Constraints, and Queries and Views

---

---

# SQL

- SQL is a comprehensive database language: It has statements for data definitions, queries, and updates.
  - It is both a DDL and a DML.
-

---

# SQL Data Definition and Data Types

- SQL uses the terms **table**, **row** and **column** for the formal relational model terms relation, tuple, and attribute, respectively.
  - The main SQL command for data definition is the CREATE statement, which can be used to create schemas, tables (relations), and domains (as well as other constructs such as views, and triggers).
-

---

# Data Definition Language

- Data definition language, which is usually part of a database management system, is used to define and manage all attributes and properties of a database, including row layouts, column definitions, key columns, file locations, and storage strategy
  - A DDL statement supports the definition or declaration of database objects such as databases, tables, and views
-

---

# Data Definition Language

Most DDL statements take the following form:

- CREATE *object\_name*
  - ALTER *object\_name*
  - DROP *object\_name*
-

---

# CREATE TABLE

USE Northwind

CREATE TABLE Importers

(

    CompanyID int NOT NULL,

    CompanyName varchar(40) NOT NULL,

    Contact varchar(40) NOT NULL

)

---

---

# ALTER TABLE

- The ALTER TABLE statement enables you to modify a table definition by altering, adding, or dropping columns and constraints or by disabling or enabling constraints and triggers.
-

---

# ALTER TABLE

- The following statement will alter the Importers table in the Northwind database by adding a column named ContactTitle to the table.

USE Northwind

ALTER TABLE Importers

ADD ContactTitle varchar(20) NULL

---



---

# DROP TABLE

- The DROP TABLE statement removes a table definition and all data, indexes, triggers, constraints, and permission specifications for that table.
  - Any view or stored procedure that references the dropped table must be explicitly dropped by using the DROP VIEW or DROP PROCEDURE statement.
-

---

# DROP TABLE

- The following statement drops the Importers table from the Northwind database.

USE Northwind

DROP TABLE Importers

---

---

# Data Control Language

- Data control language is used to control permission on database objects. Permissions are controlled by using the SQL-92 GRANT and REVOKE statements and the Transact-SQL DENY statement.
-

---

# GRANT

- The GRANT statement creates an entry in the security system that enables a user in the current database to work with data in that database or to execute specific Transact-SQL statements.
  - The following statement grants the Public role SELECT permission on the Customers table in the Northwind database:
-

---

# GRANT

USE Northwind  
GRANT SELECT  
ON Customers  
TO PUBLIC

---

---

## REVOKE

- The REVOKE statement removes a previously granted or denied permission from a user in the current database.
- The following statement revokes the SELECT permission from the Public role for the Customers table in the Northwind database:

```
USE Northwind
REVOKE SELECT
ON Customers
TO PUBLIC
```

---

---

# DENY

- The DENY statement creates an entry in the security system that denies a permission from a security account in the current database and prevents the security account from inheriting the permission through its group or role memberships.

```
USE Northwind  
DENY SELECT  
ON Customers  
TO PUBLIC
```

---

---

# Data Manipulation Language

Data manipulation language is used to select, insert, update, and delete in the objects defined with DDL.

---



---

# SELECT

The SELECT statement retrieves rows from the database and enables the selection of one or many rows or columns from one or many tables. SELECT statement, sometimes called a **mapping** or a **select-from-where block**, is formed of the three clauses SELECT, FROM, and WHERE and has the following form

```
SELECT <attribute list>  
FROM    <table list>  
WHERE   <condition>
```

---

---

# SELECT

Where

<attribute list> is a list of attribute names whose values are to be retrieved by the query.

<table list> is a list of the relation names required to process the query.

<condition> is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query.

---

---

QUERY 1: Retrieve the birth date and address of the employee(s) whose name is 'John B. Smith'.

```
SELECT      Bdate, Address  
FROM    EMPLOYEE  
WHERE Fname = 'John' AND Minit = 'B'  
AND Lname = 'Smith';
```

---

---

QUERY 2: Retrieve the name and address of all employees who work for the 'Research' department.

■ **SELECT FNAME, LNAME,  
ADDRESS  
FROM EMPLOYEE,  
DEPARTMENT  
WHERE DNAME='Research'  
AND DNUMBER=DNO**

---

---

QUERY 3: For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birthdate.

```
SELECT      PNUMBER, DNUM, LNAME,  
            BDATE, ADDRESS  
FROM  PROJECT, DEPARTMENT,  
EMPLOYEE  
WHERE DNUM=DNUMBER AND  
MGRSSN=SSN AND PLOCATION='Stafford'
```

---

---

QUERY 4: The following statement retrieves the CustomerID, CompanyName, and ContactName data for companies who have a CustomerID value equal to alfki or anatr. The result is ordered according to the ContactName value:

USE Northwind

SELECT CustomerID, CompanyName,  
ContactName

FROM Customers

WHERE (CustomerID = 'alfki' OR CustomerID  
='anatr')

ORDER BY ContactName

---

## Ambiguous Attribute Names, Aliasing, and Tuple Variables

- In SQL the same name can be used for two (or more) attributes as long as the attributes are in different relations.
- If this is the case, and a query refers to two or more attributes with the same name, we must qualify the attribute name with the relation name to prevent ambiguity.
- This is done by prefixing the relation name to the attribute name and separating the two by a period.

---

## Example: Ambiguous Attribute Names

```
SELECT Fname, EMPLOYEE.Name, Address  
FROM EMPLOYEE, DEPARTMENT  
WHERE DEPARTMENT.NAME='Research'  
AND DEPARTMENT.Dnumber= EMPLOYEE.  
.Dnumber ;
```

---



---

## Ambiguous Attribute Names, Aliasing, and Tuple Variables

- Ambiguity also arises in the case of queries that refer to the same relation twice, as in the following example.
- QUERY 5: For each employee, retrieve the employee's name, and the name of his or her immediate supervisor.

```
Q5:SELECT    E.FNAME, E.LNAME, S.FNAME,  
            S.LNAME  
      FROM    EMPLOYEE AS E, EMPLOYEE AS S  
     WHERE    E.SUPERSSN=S.SSN
```

---

---

## Ambiguous Attribute Names, Aliasing, and Tuple Variables

- In this case, we are allowed to declare alternative relation names E and S, called **aliases** or **tuple variables**, for the EMPLOYEE relation. An alias can follow the keyword **AS**, as shown in Q5,
-

---

## Unspecified WHERE Clause and Use of the Asterisk

- A missing WHERE clause indicates no condition on tuple selection; hence all tuples of the relation specified in the FROM clause qualify and are selected for the query result.
  - If more than one relation is specified in the FROM clause and there is no WHERE clause, then CROSS PRODUCT – all possible tuple combinations – of these relations is selected.
-

- 
- For example, Query 6 selects all EMPLOYEE Ssns and Query 7 selects all combinations of an EMPLOYEE Ssn and a DEPARTMENT Dname.

- Q6: **SELECT**            Ssn  
          **FROM**            EMPLOYEE

- Q7: **SELECT** Ssn, Dname  
          **FROM**    EMPLOYEE, DEPARTMENT

---

---

## Use of Asterisk

- To retrieve all the attribute values of the selected tuples, we do not have to list the attribute names explicitly in SQL; we just specify an asterisk (\*), which stands for all the attributes.

```
SELECT  *  
FROM    EMPLOYEE  
WHERE   Dn = 5
```

---

---

## Use of Asterisk

```
SELECT  *  
FROM    EMPLOYEE, DEPARTMENT  
WHERE    Dname = 'Research' AND Dn =  
          Dnumber
```

```
SELECT  *  
FROM    EMPLOYEE, DEPARTMENT
```

---

---

# Aggregate Functions in SQL

- Grouping and aggregation are required in many database applications, SQL has features that incorporate these concepts.
  - A number of built-in functions exist: **COUNT**, **SUM**, **MAX**, **MIN**, and **AVG**.
  - The COUNT function returns the number of tuples or values as specified in a query.
-

---

# Aggregate Functions in SQL

- The functions SUM, MAX, MIN, and AVG are applied to a set of multiset of numeric values and return, respectively, the sum, maximum value, minimum value, and average (mean) of those values.
  - These functions can be used in the SELECT clause or in a HAVING clause (which we introduce later).
-



---

# Aggregate Functions in SQL

- The functions MAX and MIN can also be used with attributes that have nonnumeric domains if the domain values have a total ordering among one another.
  - We illustrate the use of these functions with example queries.
-

---

Find the sum of the salaries of all employees, the maximum salary, the minimum salary, and the average salary.

- **SELECT SUM(Salary), MAX(Salary),  
MIN(Salary),AVG(Salary)  
FROM EMPLOYEE**

---

Query 9: Find the sum of the salaries of all employees of the 'Research' department, as well as the maximum salary, the minimum salary, and the average salary in this department.

■ Q9: **SELECT SUM(Salary), MAX(Salary),  
MIN(Salary),AVG(Salary)  
FROM (EMPLOYEE JOIN  
DEPARTMENT ON Dno=Dnumber)  
WHERE Dno=Dnumber AND  
Dname='Research'**

---

---

Queries 10 and 11. Retrieve the total number of employees in the company (Q10) and the number of employees in the 'Research' department (Q11).

■ Q10:      **SELECT    COUNT(\*)**  
             **FROM       EMPLOYEE**

■ Q11:      **SELECT    COUNT(\*)**  
             **FROM       EMPLOYEE, DEPARTMENT**  
             **WHERE      Dno=Dnumber AND**  
                 **Dname='Research'**

---

---

## Grouping: The **GROUP BY** and **HAVING** Clauses

- In many cases we want to apply the aggregate functions to subgroups of tuples in a relation, where the subgroups are based on some attribute values.
  - For example, we may want to find the average salary of employees in each department or the number of employees who work on each project.
-

---

## Grouping: The **GROUP BY** and **HAVING** Clauses

- In these cases we need to **partition** the relation into nonoverlapping subsets (or groups) of tuples.
  - Each group (partition) will consist of the tuples that have the same value of some attribute(s), called the **grouping attribute(s)**.
-

---

## Grouping: The **GROUP BY** and **HAVING** Clauses

- We can then apply the function to each such group independently. SQL has a **GROUP BY** clause for this purpose.
  - The **GROUP BY** clause specifies the grouping attributes, which should also appear in the **SELECT** clause, so that the value resulting from applying each aggregate function to a group of tuples appears along with the value of the grouping attribute(s).
-

---

For each department, retrieve the department number, the number of employees in the department, and their average salary.

```
SELECT Dno, COUNT(*), AVG(Salary)  
FROM EMPLOYEE  
GROUP BY Dno
```

---



---

## GROUP BY EXAMPLE

```
SELECT productid, count(*), sum(quantity)
FROM [order details]
GROUP BY productid
```

The above will group by productid, count the orders per product and return the total quantity of each product that has been ordered from the Order Details table

---

---

## GROUP BY EXAMPLE

For each project, retrieve the project number, the project name, and the number of employees who work on that project.

```
SELECT      Pnumber, Pname, COUNT(*)  
FROM        PROJECT, WORKS_ON  
WHERE        Pnumber = Pno  
GROUP BY    Pnumber, Pname;
```

---

---

## HAVING Clauses

- When we want to retrieve the values of functions for only groups that satisfy certain conditions, then we use the HAVING Clause.
  - For example, suppose that we want to modify previous so that only projects with more than two employees appear in the result.
  - SQL provides a **HAVING** clause, which can appear in conjunction with a GROUP BY clause, for this purpose.
-

---

# HAVING Clause

- HAVING provides a condition on the group of tuples associated with each value of the grouping attributes. Only that satisfy the condition are retrieved in the result of the query.
-

---

## HAVING Clause Example

Query: For each project on which more than two employees work, retrieve the project number, the project name, and the number of employees who work on the project.

```
SELECT      Pnumber, Pname, COUNT(*)  
FROM  PROJECT, WORKS_ON  
WHERE Pnumber = Pno  
GROUP BY  Pnumber, Pname  
HAVING COUNT(*) > 2
```

---

---

## The ORDER BY Clause

- The ORDER BY clause sorts a query result by one or more columns (up to 8060 bytes).
  - A sort can be ascending (ASC) or descending (DESC). If neither is specified, ASC is assumed. If more than one column is named in the ORDER BY clause, sorts are nested.
-

---

## The ORDER BY Clause - Example

USE Pubs

SELECT Pub\_id, Type, Title\_id, Price

FROM Titles

ORDER BY Pub\_id DESC, Type,  
Price

---

---

# INSERT, DELETE and UPDATE Statements in SQL

- In its simplest form, INSERT is used to add a single tuple to a relation.
  - We must specify the relation name and a list of values for the tuple.
  - The values should be listed in the same order in which the corresponding attributes were specified in the CREATE TABLE command.
-



---

## INSERT -Example

- For example, to add a new tuple to the EMPLOYEE relation ; we can use

```
INSERT INTO      EMPLOYEE
VALUES ('Richard','K','Marini', '653298653',
'30-DEC-52','98 Oak Forest,Katy,TX', 'M',
37000,'987654321', 4 )
```

---

---

## INSERT – Example 2

USE Northwind

INSERT INTO Territories

VALUES (98101, 'Seattle', 2)

---

## INSERT –Type 2

- A second form of INSERT statement allows the user to specify explicit attribute names that correspond to the values provided in the INSERT command.
- This is useful if a relation has many attributes but only a few of those attributes are assigned values in the new tuple.

---

## INSERT

- The values must include all attributes with NOT NULL specification and no default value.
  - Attributes with NULL allowed or DEFAULT values are the ones that can be left out.
-

---

## INSERT –Type 2 Example

- For example, to enter a tuple for a new EMPLOYEE for whom we know only the Fname, Lname, Dno, and Ssn attributes, we can use the statement below:

```
INSERT INTO EMPLOYEE (FNAME,  
LNAME, Dno, SSN)  
VALUES ('Richard', 'Marini',4, '653298653')
```

---

---

## UPDATE

- The UPDATE command is used to modify attribute values of one or more selected tuples.
  - As in the DELETE command, a WHERE clause in the UPDATE command selects the tuples to be modified from a single relation.
-

---

# UPDATE

- However, updating a primary key value may propagate to the foreign key values of tuples in other relations if such a referential triggered action is specified in the referential integrity constraints of the DDL.
  - An additional **SET** clause in the UPDATE command specifies the attributes to be modified and their new values.
-

---

## UPDATE - Example

- For example, to change the location and controlling department number of project number 10 to 'Bellaire' and 5 respectively, we use the statement below:

```
UPDATE      PROJECT
SET          PLOCATION = 'Bellaire', DNUM = 5
WHERE       PNUMBER=10
```





---

# DELETE STATEMENT

The DELETE statement removes rows from a table

DELETE

FROM [table]

WHERE {condition}

---

---

# DELETE STATEMENT - Example

This example deletes all rows from the **authors** table.

USE pubs

DELETE authors

---

---

## DELETE STATEMENT - Example

- This example deletes all rows in which **au\_lname** is McBadden.

USE pubs

DELETE FROM authors

WHERE au\_lname = 'McBadden'

---

---

# VIEWS

- A **view** in SQL terminology is a single table that is derived from other tables.
  - These other tables can be base tables or previously defined views.
  - A view does not necessarily exist in physical form; it is considered a **virtual table**, in contrast to base tables, whose tuples are actually stored in the database.
-

---

## VIEWS

- We may frequently issue queries that retrieve the employee name and the project names that the employee works on.
  - Rather than having to specify the join of the EMPLOYEE, WORKS\_ON, and PROJECT tables every time we issue that query, we can define a view that is a result of these joins.
-

---

## VIEWS

- Then we can issue queries on the view, which are specified as single table retrievals rather than as retrievals involving two joins on three tables.
  - We call the EMPLOYEE, WORKS\_ON, and PROJECT tables the **defining tables** of the view.
-

---

## VIEW – Example

```
CREATE VIEW WORKS_ON1  
AS SELECT    Fname, Lname, Pname, Hours  
FROM        EMPLOYEE, PROJECT, WORKS_ON  
WHERE        Ssn=Essn AND Pno =Pnumber
```

```
CREATE VIEW DEPT_INFO(Dept_name, No_of_emps, Total_sal)  
AS SELECT    Dname, COUNT (*), SUM(Salary)  
             FROM    DEPARTMENT,EMPLOYEE  
             WHERE    Dnumber=Dno  
             GROUP BY Dnumber
```

---

---

## Querying A View

- we can utilise the WORKS\_ON1 view and specify the query as below:

```
SELECT  Fname, Lname  
FROM    WORKS_ON1  
WHERE    Pname = 'ProjectX'
```



---

## Using Advanced Query Techniques To Access Data

- One of these techniques is to combine the contents of two or more tables to produce a result set that incorporates rows and columns from each table.
  - Another technique is to use subqueries, which are SELECT statements nested inside other SELECT, INSERT, UPDATE, or DELETE statements.
  - Subqueries can also be nested in other subqueries.
-

# Using Joins To Retrieve Data

- By using joins, you can retrieve data from two or more tables based on logical relationships between the tables. Joins can be specified in either the FROM or WHERE clauses.
- The join conditions combine with the WHERE and HAVING search conditions to control the rows that are selected from the base tables referenced in the FROM clause.
- Specifying the join conditions in the FROM clause, however, helps separate them from any other search conditions that might be specified in a WHERE clause and is the recommended method for specifying joins.

---

# JOINS

- When multiple tables are referenced in a single query, all column references must be unambiguous.
  - The table name must be used to qualify any column name that is duplicated in two or more tables referenced in a single query.
-

---

# JOINS

- The select list for a join can reference all of the columns in the joined tables or any subset of the columns.
  - The select list is not required to contain columns from every table in the join.
-

# JOINS

- For example, in a three-table join, only one table can be used as a bridge from one of the other tables to the third table, and none of the columns from the middle table have to be referenced in the select list.
- Although join conditions usually use the equals sign ( = ) comparison operator, other comparison or relational operators can be specified ( as can other predicates). Most joins can be categorized as inner joins or other joins.

---

# JOINS

- Inner joins return rows only when there is at least one row from both tables that matches the join condition, eliminating the rows that do not match with a row from the other table.
  - Outer joins, however, return all rows from at least one of the tables or views mentioned in the FROM clause as long as these rows meet any WHERE or HAVING search conditions.
-

---

# INNER JOIN

An inner join is a join in which the values in the columns being joined are compared through the use of a comparison operator.

- The following SELECT statement uses an inner join to retrieve data from the Publishers table and the Titles table in the Pubs database:

```
SELECT t.Title, p.Pub_name  
FROM Publishers AS p INNER JOIN Titles AS t  
ON p.Pub_id = t.Pub_id  
ORDER BY Title ASC
```

---

# Outer Joins

- There are three types of outer joins: left, right, and full.
- All rows retrieved from the left table are referenced with a left outer join, and all rows from the right table are referenced in a right outer join.
- All rows from both tables are returned in a full outer join.



## Using Left Outer Joins

- A result set generated by a SELECT statement that includes a left outer join includes all rows from the table referenced to the left of LEFT OUTER JOIN.
- The only rows that are retrieved from the table to the right are those that meet join condition.

---

## Using Left Outer Joins

- In the following SELECT statement, a left outer join is used to retrieve the authors' first names, and (when applicable) the names of any publishers that are located in the same cities as the authors:

---

---

## Using Left Outer Joins

USE Pubs

```
SELECT a.Au_fname, a.Au_lname, p.Pub_name  
FROM   Authors a LEFT OUTER JOIN Publishers p  
ON a.City = p.City  
ORDER BY p.Pub_name ASC, a.Au_lname ASC,  
a.Au_fname ASC
```

---

---

## Using Right Outer Joins

- A result set generated by a SELECT statement that includes a right outer join includes all rows from the table referenced to the right of RIGHT OUTER JOIN.
  - The only rows that are retrieved from the table to the left are those that meet the join condition.
  - In the following SELECT statement, a right outer join is used to retrieve the list of publishers and authors' first names, and last names, if those authors are located in the same cities as the publishers:
-

---

## Using Right Outer Joins

USE Pubs

SELECT a.Au\_fname, a.Au\_lname,  
p.Pub\_name

FROM Authors a RIGHT OUTER JOIN  
Publishers p

ON a.City = p.City

ORDER BY p.Pub\_name ASC, a.Au\_lname  
ASC, a.Au\_fname ASC

---

---

# Using Full Outer Joins

- A result set generated by a SELECT statement that includes a full outer join includes all rows from both tables, regardless of whether the tables have a matching value (as defined in the join condition).
  - In the following SELECT statement, a full outer join is used to retrieve the list of publishers and authors' first and last names:
-

---

## Using Full Outer Joins

USE Pubs

```
SELECT a.Au_fname, a.Au_lname, p.Pub_name  
FROM   Authors a FULL OUTER JOIN Publishers p  
ON a.City = p.City  
ORDER BY p.Pub_name ASC, a.Au_lname ASC,  
         a.Au_fname ASC
```

---

## Defining Subqueries inside SELECT Statement

- A subquery is a SELECT statement that returns a single value and is nested inside a SELECT, INSERT, UPDATE or DELETE statement or inside another subquery.
- A subquery can be used anywhere an expression is allowed. A subquery is also called an inner query or inner select, while the statement containing a subquery is called an outer select.
- In the following example, a subquery is nested in the WHERE clause of the outer SELECT statement:



---

## Subquery - Example

```
USE Northwind
SELECT ProductName
FROM Products
WHERE UnitPrice =
    {
        SELECT UnitPrice
        FROM Products
        WHERE ProductName = 'Sir Rodney' 's
                                Scones'
    }
```

---

## Subqueries

- If a table only appears in a subquery and not in the outer query, then columns from that table cannot be included in the output (the select list of the outer query).

## Types of Subqueries

Subqueries can be specified in many places within a SELECT statement. Statements that include a subquery usually take one of the following formats:

WHERE <expression> [NOT] IN (<subquery>)

WHERE <expression> <comparison\_operator>  
[ANY | ALL] (<subquery>)

WHERE [NOT] EXISTS (<subquery>)

---

## Subqueries that Are used with IN or NOT IN

- The result of a subquery introduced with IN (or with NOT IN) is a list of zero or more values.
  - After the subquery returns the result, the outer query makes use of it.
  - In the following example, a subquery is nested inside the WHERE clause, and the IN keyword is used:
-

---

## Subquery - Example

```
USE Pubs
SELECT Pub_name
FROM Publishers
WHERE Pub_id IN
(
  SELECT Pub_id
  FROM Titles
  WHERE Type = 'business'
)
```

---

---

## Subquery

- Subqueries introduced with the NOT IN keywords also return a list of zero or more values.
  - The query is exactly the same as the one in subqueries with IN, except that NOT IN is substituted for IN.
-

---

## Subqueries that Are Used with Comparison Operators

- Comparison operators that introduce a subquery can be modified with the keyword ALL or ANY.
  - Subqueries introduced with a modified comparison operator return a list of zero or more values and can include a GROUP BY or HAVING clause.
  - These subqueries can be restated with EXISTS.
-

---

## Subqueries that Are Used with Comparison Operators

- The ALL and ANY keywords each compare a scalar value with a single-column set of values.
  - The ALL keyword applies to every value, and the ANY keyword applies to at least one value.
  - In the following example, the greater than (>) comparison operator is used with the ANY keyword:
-

---

## Subquery - Example

USE Pubs

SELECT Title

FROM Titles

WHERE Advance > ANY

(

    SELECT Advance

    FROM Publishers INNER JOIN Titles

    ON Titles.Pub\_id = Publishers.Pub\_id

    AND Pub\_name = 'Algodata Infosystems'

)

---



---

## Subqueries that Are Used with EXISTS and NOT EXISTS

- When a subquery is introduced with the keyword EXISTS, it functions as an existence test.
  - The WHERE clause of the outer query tests for the existence of rows returned by the subquery.
  - The subquery does not actually produce any data; instead, it returns a value of TRUE or FALSE.
  - In the following example, the WHERE clause in the outer SELECT statement contains the subquery and uses the EXISTS keyword.
-

---

## Subquery - Example

USE Pubs

SELECT Pub\_name

FROM Publishers

WHERE EXISTS

(

SELECT \* FROM Titles

WHERE Titles.Pub\_id = Publishers.Pub\_id

AND Type = 'Business'

)

---

# Implementing Integrity Constraints

- A constraint is a property assigned to a table or a column within a table that prevents invalid data values from being placed in specified column(s).
- Examples include UNIQUE, constraints, PRIMARY KEY constraints and FOREIGN KEY constraints

---

## Creating PRIMARY KEY Constraints

You can create a PRIMARY KEY onstraint by using one of the following methods:

- Creating the constraint when the table is created (as part of the table definition)
  - Adding the constraint to an existing table, provided that no other PRIMARY KEY constraint already exists.
-

---

## PRIMARY KEY -EXAMPLE

```
CREATE TABLE Table1  
(  
    col1 int PRIMARY KEY,  
    col2 varchar(30)  
)
```

---

---

## PRIMARY KEY – table level constraint

```
CREATE TABLE Table1  
(  
    col1 int ,  
    col2 varchar(30),  
    CONSTRAINT table_pk PRIMARY KEY (col1)  
)
```

---

---

## Composite Key

```
CREATE TABLE FactoryProcess  
(  
    EventType INT,  
    EventTime DATETIME,  
    EventSite CHAR(50),  
    EventDesc CHAR(1024),  
    CONSTRAINT event_key PRIMARY KEY  
    (EventType, EventTime)  
)
```

---

---

# Adding a Primary to an Existing Table

```
ALTER TABLE Table1
```

```
ADD CONSTRAINT table_pk PRIMARY KEY (col1)
```





---

# UNIQUE CONSTRAINTS

- You can use UNIQUE constraints to ensure that no duplicate values are entered in specific columns that do not participate in a primary key.
  - Although both a UNIQUE constraint and a PRIMARY KEY constraint enforce uniqueness, you should use a UNIQUE constraint instead of a PRIMARY KEY constraint in the following situations:
-

---

# UNIQUE CONSTRAINTS

- **If a column (or combination of columns) is not the primary key.** Multiple UNIQUE constraints can be defined on a table, whereas only one PRIMARY KEY constraint can be defined on a table.
-

---

# UNIQUE CONSTRAINTS

- **If a column allows null values.** UNIQUE constraints can be defined for columns that allow null values, whereas PRIMARY KEY constraints can be defined only on columns that do not allow null values.
  - A UNIQUE constraint can also be referenced by a FOREIGN KEY constraint.
-

---

# Creating UNIQUE Constraints

- You can create a UNIQUE constraint in the same way that you create a PRIMARY KEY constraint:
  - By creating the constraint when the table is created (as part of the table definition)
  - By adding the constraint to an existing table, provided that the column or combination of columns comprising the UNIQUE constraint contains only unique or NULL values. A table can contain multiple UNIQUE constraints.
-

---

# Creating UNIQUE Constraints

- You can use the same Transact-SQL statements to create a UNIQUE constraint that you used to create a PRIMARY KEY constraint.
  - Simply replace the words PRIMARY KEY with the word UNIQUE.
-

# FOREIGN KEY Constraints

- A foreign key is a column or combination of columns used to establish and enforce a link between the data in two tables.
- Create a link between two tables by adding a column (or columns) to one of the tables and defining those columns with a FOREIGN KEY constraint.
- The columns will hold the primary key values from the second table. A table can contain multiple FOREIGN KEY constraints.

---

# Creating FOREIGN KEY Constraints

You can create a FOREIGN KEY constraint by using one of the following methods:

- Creating the constraint when the table is created (as part of the table definition)
  - Adding the constraint to an existing table, provided that the FOREIGN KEY constraint is linked to an existing PRIMARY KEY constraint or a UNIQUE constraint in another (or the same) table
-

---

# Foreign Key Constraints

```
CREATE TABLE Table1  
(  
  Col1 INT PRIMARY KEY,  
  Col2 INT REFERENCES Employees(EmployeeID)  
)
```

---



---

## table-level FOREIGN KEY constraint

```
CREATE TABLE Table1
(
  Col1 INT PRIMARY KEY,
  Col2 INT,
  CONSTRAINT col2_fk FOREIGN KEY (Col2)
  REFERENCES Employees (EmployeeID)
)
```

---

---

# ALTER TABLE statement to add a FOREIGN KEY

```
ALTER TABLE Table1  
ADD CONSTRAINT col2_fk FOREIGN KEY  
(Col2) REFERENCES Employees (EmployeeID)
```

---

